

CRED Platform: Data Pipeline Design Doc

Author: Richa Patel

Date: October 2025

1. Goal

Design a scalable, reliable, and cost-effective data platform to support AI-driven analytics and dashboards for high-volume user event data.

Requirements:

- Handle billions of rows efficiently
- Serve both analytics and ML workloads
- Ensure data quality, reproducibility, and incremental processing
- Balance engineering effort with Series A startup resource constraints

2. Event Data Overview

Sample event:

```
{  
  "user_id": "abc123",  
  "event_type": "purchase",  
  "event_timestamp": "2025-09-01T12:34:56Z",  
  "metadata": {  
    "amount": 42.50,  
    "currency": "USD"  
  }  
}
```

Events originate from web and mobile apps and may arrive in real-time or batch.

3. Pipeline Architecture

Flow Overview:

CRED Data Platform: Event Pipeline Overview

1. Raw Event Ingestion

- Events are collected from web and mobile apps.
- Can be ingested via GCP Pub/Sub for streaming or CSV uploads for batch.

2. Raw Table in BigQuery (events_raw)

- All incoming events are stored in their raw form.
- Maintains an immutable, auditable history of all events.

3. Staging Layer (stg_events) via dbt

- Performs data cleaning, such as removing duplicates and correcting types.
- Flattens JSON fields (e.g., metadata.amount) for easier querying.
- Validates data types and schema consistency.

4. Fact / Curated Layer (fact_user_activity) via dbt

- Aggregates key metrics per user for dashboards and ML:
 - first_event – timestamp of first activity
 - last_event – timestamp of last activity
 - total_events – number of events per user
 - total_amount – sum of purchase amounts
- Supports incremental loading to process only new data.
- Tables are partitioned by event_date for performance and cost efficiency.

5. Serving Layer: GraphQL API / ML Feature Store

- Provides curated metrics to dashboards and AI models.
- Supports daily refresh or near-real-time serving depending on business need.

3.1 Ingestion

- Choice: GCP Pub/Sub for streaming; CSV uploads for batch testing.
- Reasoning:
 - Streaming ensures near-real-time analytics for dashboards and ML features.
 - Batch ingestion provides a simple, low-cost method for testing or small data loads.
- Trade-offs:
 - Streaming introduces complexity (Pub/Sub + Dataflow).
 - Batch is simpler but slower for real-time analytics.

3.2 Storage Strategy

- Raw Layer: events_raw
 - Immutable, captures all events

- Partitioned by ingestion date for cost efficiency
- Curated Layer: fact_user_activity
 - Aggregated metrics for analytics & ML
 - Partitioned by event_date
 - Clustered by user_id for efficient lookups

Trade-offs:

- Raw layer preserves full history; good for audits and backfills
- Curated layer is optimized for query performance

3.3 Transformations (dbt)

- Staging Layer (stg_events):
 - Flatten nested JSON fields
 - Type validation & null checks
 - Standardize timestamps and currency
- Fact Layer (fact_user_activity):
 - Aggregates per user per day:
 - first_event / last_event timestamps
 - total_events / total_amount
 - Incremental logic: only process new events
 - Partitioned & clustered for performance

Trade-offs:

- dbt simplifies transformations and testing
- Incremental models reduce cost and compute time

3.4 Serving Layer

- GraphQL API or ML Feature Store
- Supports dashboards and AI model consumption
- Refresh strategy: daily incremental updates (or near real-time for critical features)

4. Scalability & Reliability

- Partitioning & Clustering: reduces query scan cost
- Incremental loading: avoids full-table recomputation

- Data quality checks:
 - Row counts, null checks, uniqueness
 - dbt tests integrated in CI/CD pipeline
- Error handling & monitoring: alert on schema changes or failed incremental runs

5. Series A Prioritization

- First priorities:
 1. Reliable raw ingestion & staging
 2. Fact table with core metrics (first_event, last_event, total_events, total_amount)
 3. Data quality tests to ensure trust in metrics
- Later priorities:
 - Near real-time dashboards
 - Feature store for ML
 - Lineage tracking & cost optimization

6. Summary

This design provides a modular, scalable, and reliable pipeline:

- Raw → Staging → Fact → Serving layers
- Supports AI/ML workloads with incremental aggregation and partitioning
- Enforces data quality via dbt tests
- Balances engineering effort with Series A resource constraints